


Machine Learning



Decision Tree

$$\text{Entropy} = -\sum p_i \log_2 p_i$$

↳ Impurity (chaos)




$0.29 = \text{Information Gain}$






Salary	Age
10000	M.A → 0
9500	Senior → 1 ←
20000	Youth → <u>2</u> ←
...	

Salary
 9500
 9700
 10000
 ... ←




X test



✓ test


Decision Tree

- One of the widely used and practical methods for classification
- Utilizes a tree structure to model relationships among the features and the potential outcomes (target attribute)
- Decision trees consist of nodes and branches.



Decision Tree

- Nodes are decision points
- Branches are the result of the decision function.
- The nodes are of three types:
 - Root Node (representing the original data) and a decision function
 - Branch Node (representing a decision function)
 - Leaf Node (holds the result of all the previous functions that flow to it)




Decision Tree

- Goal of a decision tree is to classify or predict an outcome based on a set of predictors
- For example: to Predict whether a customer will buy a product or not. Predictors: age of customer, credit rating etc
- Tree creation splits data into subsets and subsets into further smaller subsets.
- The algorithm stops splitting data when data within the subsets are sufficiently homogenous or some other stopping criterion is met

Decision Tree

- Decision Tree – visualize the increasing depth of tree




<https://share.cocalc.com/share/8b892baf91f98d0cf6172b872c8ad6694d0f7204/PythonDataScienceHandbook/notebooks/05.08-Random-Forests.ipynb?viewer=share>

Decision Tree

- After executing all the decision functions from Root Node to Leaf Node, the class of a data point is decided by the leaf node to which it reaches
- The leaf node may not contain all data points of same class. The Leaf Node belongs to the majority class
- Once a decision tree has been constructed, it is a simple matter to convert it into an equivalent set of rules that can be applied to a new data point and predict its class

Decision Tree Training

- Decision tree algorithm learns through the measure of impurity of data in a node
- Which of the following node has the most impurity?




- Impurity at a node is measured based on mixture of different classes in the target column of a node
- The objective is to minimize the impurity as much as possible at the leaf nodes

Measuring Impurity

- **Entropy**


- A box contains 6 red and 4 black balls.
- (Imagine that these represent data points of two different classes)




- Entropy of the box is calculated as : $Entropy = \sum_{i=1}^C -p_i * \log_2(p_i)$
- Entropy = - (0.6 * $\log_2(0.6)$) - (0.4 * $\log_2(0.4)$) = 0.9709506
- If we remove all red balls from the bag and then entropy will be
- Entropy = - 1.0 * $\log_2(1.0)$ - 0.0 * $\log_2(0)$ = 0
- What do you think is the interpretation of Entropy = 0?

Entropy

- Entropy ranges from 0 to 1.
- Entropy 0 means 100% information
- Entropy 1 mean maximum uncertainty
- Entropy values for a two-class variable are as follows



Entropy – Information Gain



$$I_H(D_p) = -(0.5 \log_2(0.5) + 0.5 \log_2(0.5)) = 1$$


$$-\left(\frac{3}{4} \log_2\left(\frac{3}{4}\right) + \frac{1}{4} \log_2\left(\frac{1}{4}\right)\right) = 0.81$$

$$-\left(\frac{1}{4} \log_2\left(\frac{1}{4}\right) + \frac{3}{4} \log_2\left(\frac{3}{4}\right)\right) = 0.81$$

• Information Gain = $1 - \frac{4}{8}0.81 - \frac{4}{8}0.81 = 0.19$

Entropy - example

- Decision tree to predict whether a loan given to a person would result in a write-off or not. Data consists of 30 instances. 16 belong to the write-off class and the other 14 belong to the non-write-off class.



$$E(\text{Parent}) = -\frac{16}{30} \log_2 \left(\frac{16}{30} \right) - \frac{14}{30} \log_2 \left(\frac{14}{30} \right) \approx 0.99$$

$$E(\text{Balance} < 50K) = -\frac{12}{13} \log_2 \left(\frac{12}{13} \right) - \frac{1}{13} \log_2 \left(\frac{1}{13} \right) \approx 0.39$$

$$E(\text{Balance} > 50K) = -\frac{4}{17} \log_2 \left(\frac{4}{17} \right) - \frac{13}{17} \log_2 \left(\frac{13}{17} \right) \approx 0.79$$

Weighted Average of entropy for each node:

$$\begin{aligned} E(\text{Balance}) &= \frac{13}{30} \times 0.39 + \frac{17}{30} \times 0.79 \\ &= 0.62 \end{aligned}$$

Information Gain:

$$\begin{aligned} IG(\text{Parent}, \text{Balance}) &= E(\text{Parent}) - E(\text{Balance}) \\ &= 0.99 - 0.62 \\ &= 0.37 \end{aligned}$$


Gini Index

- Gini index is another measure which gives similar results as Entropy
- It is calculated by subtracting the sum of the squared probabilities of each class from one

$$Gini = 1 - \sum_{i=1}^C (p_i)^2$$

- When uncertainty is highest, i.e. when data is evenly distributed in a node, value will be 0.5
- In perfectly classified node, values will be 0

Gini Index – Information Gain



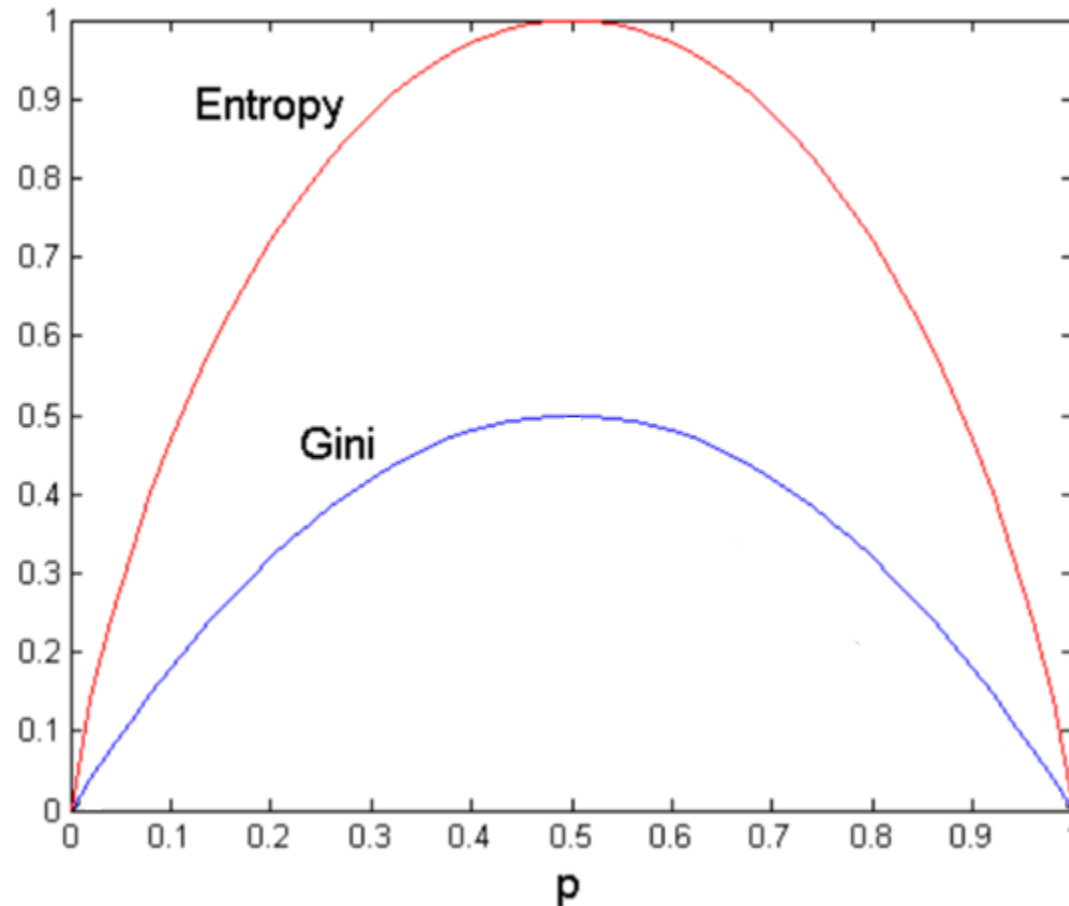
$$1 - \left(\left(\frac{3}{4} \right)^2 + \left(\frac{1}{4} \right)^2 \right) = \frac{3}{8} = 0.375$$

$$1 - \left(\left(\frac{1}{4} \right)^2 + \left(\frac{3}{4} \right)^2 \right) = \frac{3}{8} = 0.375$$

$$\text{Information Gain} = 0.5 - \frac{4}{8} 0.375 - \frac{4}{8} 0.375 = 0.125$$

Entropy and Gini Index

- Comparison for 2-class problem



Decision Tree Algorithms


- ID3 (Iterative Dichotomiser 3) was developed in 1986 by Ross Quinlan. The algorithm creates a multiway tree, finding for each node (i.e. in a greedy manner) the categorical feature that will yield the largest information gain for categorical targets
- C4.5 is the successor to ID3 and removed the restriction that features must be categorical by dynamically defining a discrete attribute
- C5.0 is latest version release under a **proprietary license**. It uses less memory and builds smaller rulesets than C4.5 while being more accurate
- CART (Classification and Regression Trees) is very similar to C4.5, but it differs in that it supports numerical target variables. It creates binary tree. Available in Scikit-learn

Decision Tree

- Advantages :
 - Simple to understand and to interpret
 - Trees can be visualized
 - Able to handle both numerical and categorical data.
 - Requires less data preparation. Does well with missing data
 - Does not assume relationship between features and target variables
- Disadvantages:
 - Decision trees can be unstable and tend to overfit. (Mitigation: Use decision trees within an ensemble)
 - Decision tree learners create biased trees if some classes dominate
 - Divide feature space in axis parallel boundaries which may not be optimum

Decision Tree - Parameters

- `max_depth` – Is the maximum length of a path from root to leaf (in terms of number of decision points). The leaf node is not split further. It could lead to a tree with leaf node containing many observations on one side of the tree, whereas on the other side, nodes containing much less observations get further split
- `min_sample_split` - A limit to stop further splitting of nodes when the number of observations in the node is lower than this value
- `min_sample_leaf` – Minimum number of samples a leaf node must have. When a leaf contains too few observations, further splitting will result in overfitting (modeling of noise in the data)
- `max_leaf_nodes` – maximum number of leaf nodes in a tree



Max Dep = 2

Average Score

(I) Train 1, 2 Test 3

(II) Train 1, 3 Test 2

(III) Train 2, 3 Test 1

Max Dep = 3

—
—
—

—
—
—

Max Dep = 4

—
—
—

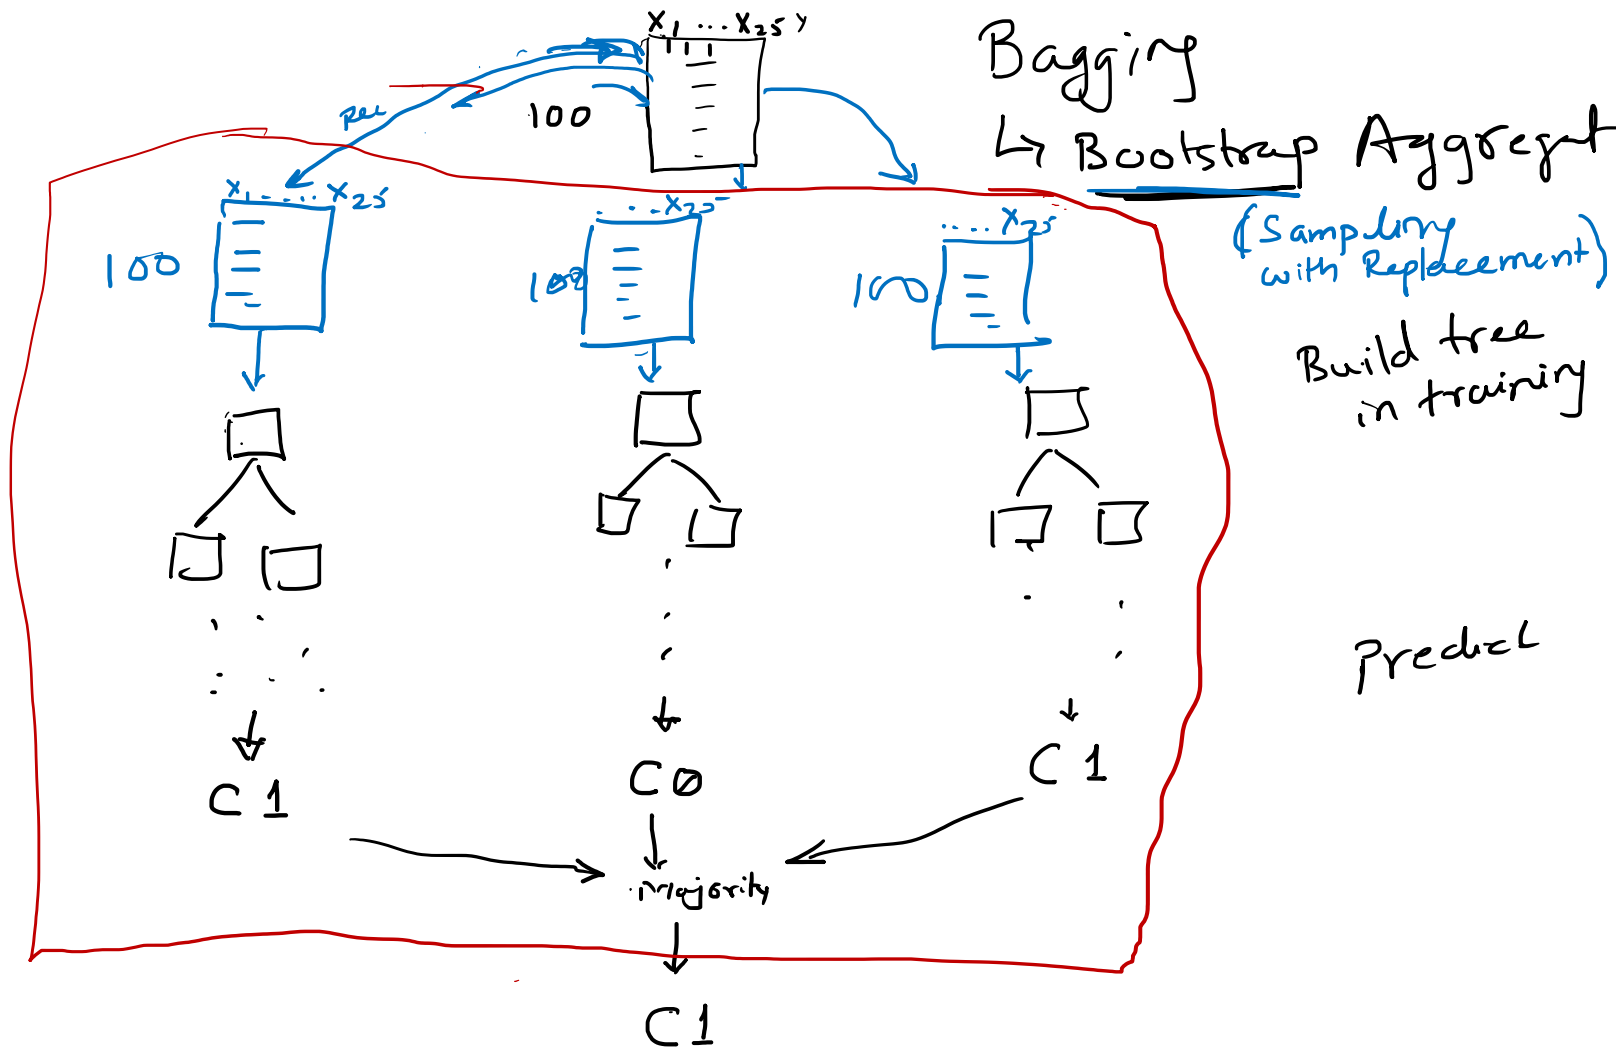
—
—
—

Max Dep = 17


—
—
—

Machine Learning




Ensemble Techniques



Random Forest




→ Bootstrap Sampling.



n column
 \sqrt{n} column.



Random
5 columns
FOR EVALUATION



TRAINING
-FIT

Random
5 column

STATQUEST.



Ensemble - Background

- What is Ensemble?
 - Do not predict using a single classifier but learn a set of classifiers
 - An ensemble of classifiers is created by combining predictions of multiple classifiers for improving prediction performance
- Why Ensemble?
 - Combining the outputs of several classifiers may reduce the risk of selecting a poorly performing classifier
 - The errors made while classifying instances by one classifier are generally averaged out by the correct classification of another classifier, so that the overall classification accuracy is improved

Ensemble - Background

- In some parts of the feature space, the different classifiers produce similar results
- In regions where the data points from different classes overlap, the classifiers give different results
- By using information from multiple classifiers, the result may be better than an individual classifier

Ensemble

- For Ensemble to work successfully, we need to ensure each learner (classifier, in this case) is slightly different
- How to achieve this?
 - Provide different data to different classifiers
 - Perform random sampling with replacement of rows
 - Provide different features as input to different classifiers
 - By adjusting the weights assigned to each data point to force an instance to focus on certain data points more

Bagging

- Bagging term is made from - **B**ootstrap **A**ggregation
- It is used to reduce the variance of a decision tree
- It creates several subsets of data from training sample chosen randomly with replacement.
- Each subset data is used to train a decision tree
- As a result, we have an ensemble of different models.
- For classification, bagging is used with voting to decide the class of an input while for regression average or median values are calculated
- This concept can be extended for other models too, but is commonly used for Decision Tree

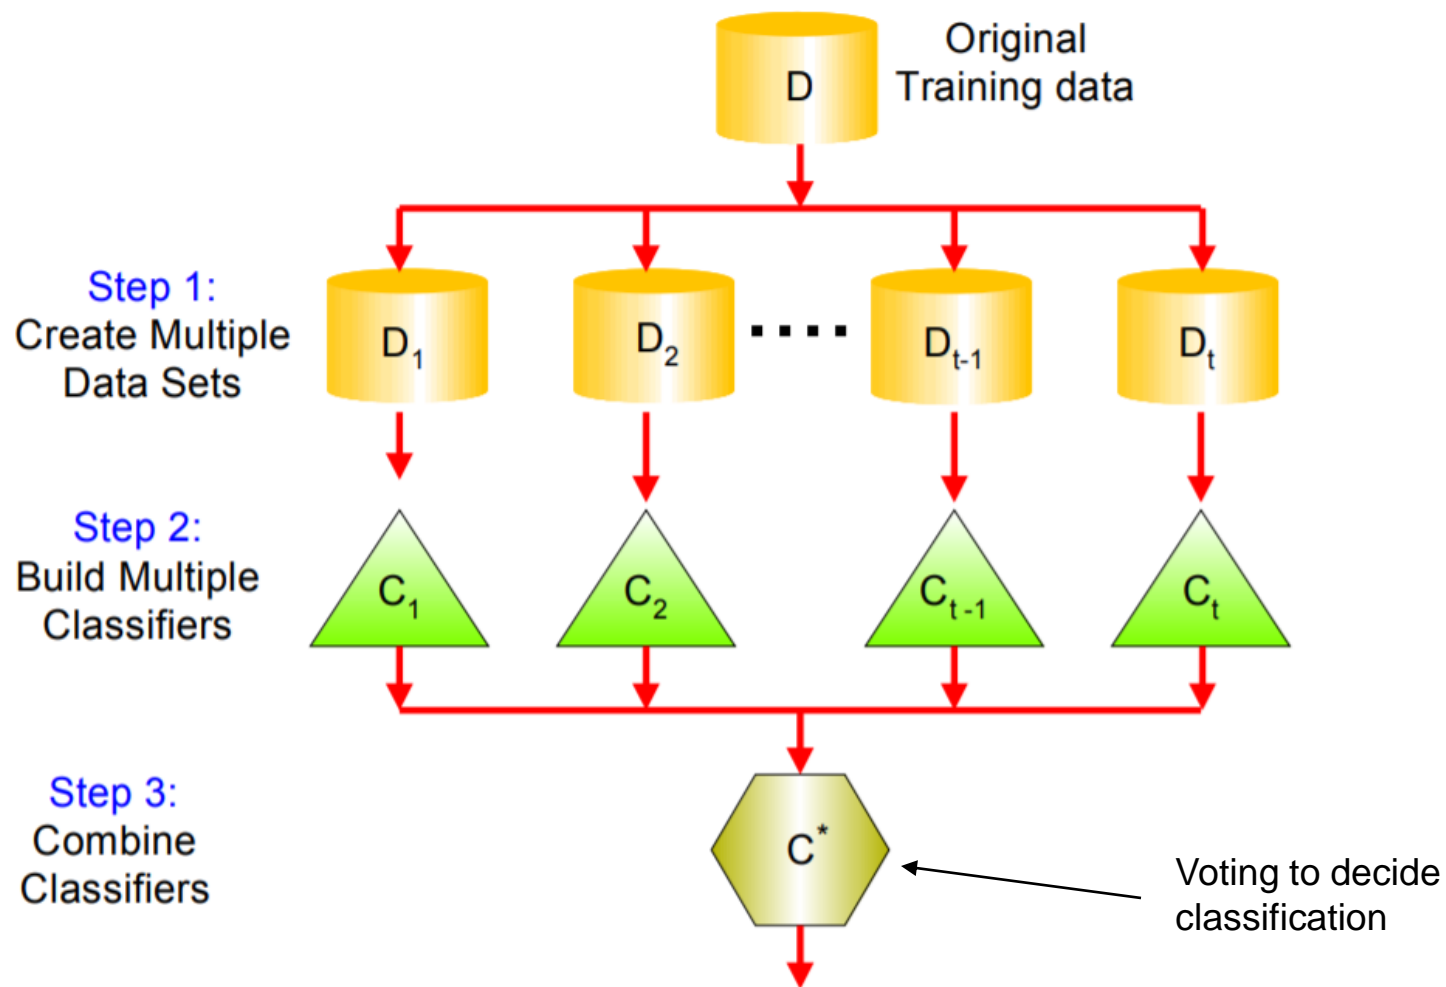
Bootstrap sampling

- Generate new training sets using sampling with replacement. It is called bootstrap sampling

Original Data	1	2	3	4	5	6	7	8	9	10
Bagging (Round 1)	7	8	10	8	2	5	10	10	5	9
Bagging (Round 2)	1	4	9	1	2	3	2	7	3	2
Bagging (Round 3)	1	8	5	10	5	5	9	6	3	7

- Some data may appear in more than one set
- Some data will appear more than once in a set

Bagging



Random Forest

- Build each tree using a sample drawn with replacement (bootstrap) from the training set
- When splitting a node during the construction of a tree, the split that is chosen is no longer the best split among all the features
- Instead, the split is picked is the best split among a random subset (say, k number of subsets) of the features
- As a result of this randomness, the bias of the forest usually slightly increases (with respect to the bias of a single non-random tree)
- Due to averaging, its variance decreases, usually more than compensating the increase in bias, hence yielding overall a better result

Random Forest

- Advantages
 - The process of averaging or combining the results of different decision trees helps to overcome the problem of overfitting
 - Less variance than a single decision tree
 - Higher accuracy than a single decision tree
- Disadvantages
 - The advantage of Decision Tree, interpretability, is lost. No interpretability
 - Need to choose the number of trees

Exercise

- Ensemble Techniques


Machine Learning

Bias Variance

Bias Variance Errors

- Bias
 - Caused by our selection of the attributes and our interpretation of their influence on each other
 - The real model in the universe / population may have many more attributes and the attributes interacting in different ways not reflected in our model
- Variance
 - Different test data – gives very different scores. Variance is the amount that the estimate of the target function will change if different training data was used
 - Caused by overfitting of model

Bias – Variance trade-off



- Select the right complexity model to trade-off Bias and Variance errors